

R Programming

Scoping

Variable Scope

- A *variable* is pairing of a name and a value.
- Variables can be created by assignment. The assignment

`x = 10`

indicates that the value “10” and the name “x” are to be paired.

- R has both *local* and *global* variables.
- Global variables are created by assignments at top level.
- Local variables are created by assigning values within functions.

Example: Same Name, Different Variables

There are two “x” variables in the following code.

```
x = 10                                # <-- global

fun =
  function() {
    x = 20                            # <-- local to fun
  }
```

Inside “fun,” the value of “x” is “20.”

Outside “fun,” the value of “x” is “10.”

Changes made to one “x” do not affect the other.

Visibility of Global Variables

Inside a function, the values of both local and global variables are visible.

```
> y = 20
> fun =
      function() {
          x = 30
          x + y
      }
> fun()
[1] 50
```

Here, “**fun**” combines the value of a local variable, “**x**,” and the value of a global variable, “**y**.”

Nested Functions

When R functions are nested, the variables of the outer functions are visible within the nested functions.

```
> outer =  
  function() {  
    inner =  
      function() {  
        y = 20      # y is local to inner  
        x + y  
      }  
    x = 10           # x is local to outer  
    inner()  
  }  
> outer()  
[1] 30
```

Scoping Rules

- The scoping rules of a (computer) language govern how the value of variables can be determined.
- The scope of a variable is the region of code where that variable has meaning.
- In R, a local variable has meaning only within the function it is local to (including any functions that are nested within that function).
- When a value is sought for a given name, the local scope is inspected for a value, then any nesting scopes (i.e. functions) and finally a global value is sought.

Function Parameters, Arguments and Variables

- Assignment is not the only way that variables are created.
- The association of function's arguments with the function's formal parameters also creates variables.
- In this case the variable name is the formal parameter name and the (initial) value of the variable is corresponding function argument.

Example: Parameters and Arguments

- In the function definition

```
fun =  
    function(a, b)  
        a^2 + b^2
```

the *formal parameters* of the function are “a” and “b.”

- In the function call

```
fun(10, 20)
```

the *arguments* to the function are “10” and “20.”

- The arguments and parameters are paired as variables.

A Simple Example

The following function adds a local variable “u” to a global variable “x.”

```
> add.x.to = function(u) x + u
```

```
> x = 10
```

```
> add.x.to(20)
```

```
[1] 30
```

```
> x = 20
```

```
> add.x.to(20)
```

```
[1] 40
```

More Complexity

Now let's embed the “`add.x.to`” function inside another function.

```
> add =  
  function(x, y) {  
    add.x.to = function(u) x + u  
    add.x.to(y)  
  }
```

```
> add(1, 2)  
[1] 3
```

Now “`x`” is no longer global. It is a variable that is local to the function “`add`.”

An Important Change

Suppose now we change the embedding function as follows.

```
> make.add.to =  
  function(x, y) {  
    add.x.to = function(u) x + u  
    add.x.to  
  }
```

```
> add1 = make.add.to(1)
```

Now, rather than returning the value that results from applying “`add.x.to`” to a value, we return the function itself

(The variable “`y`” is no longer used and we can drop from the definition of “`add.x.to`.”)

What Does The Function Do?

We can now try applying the function “`add1`” to various arguments.

```
> add1(1)
[1] 2
> add1(2)
[1] 3
> add1(10)
[1] 11
```

Clearly “`add1`” is a function that adds “1” to its argument.

Does This Always Work?

Let's try making other functions with “`make.add.to.`”

```
> add10 = make.add.to(10)
```

```
> add10(1)
```

```
[1] 11
```

```
> add10(100)
```

```
[1] 110
```

Clearly this is something that works in general!

Why?

The Process Explained

The “`make.add.to`” function is defined as follows.

```
> make.add.to =  
    function(x)  
      function(u) x + u
```

When “`make.add.to`” is called, a variable called “`x`” is created and set to the value of the argument to “`make.add.to`”.

The function defined inside “`make.add.to`” has access to the variable “`x`.”

This remains true, even after “`make.add.to`” has returned.

Closures

- The function returned by “`make.add.to`” is said to *close over* the value of “`x`.”
- The function retains access to the variable “`x`” and can use its value whenever it is called.
- Functions that do this are called *closures*.
- Forming closures is one of R’s most important capabilities.
- It is also one of the least-understood and least-used.

Defining A Likelihood

The following function creates closures that represent Poisson likelihood functions.

```
> mk.pois.lk =  
  function(x)  
    function(lambda) prod(dpois(x, lambda))
```

Given a vector containing the values of n Poisson random variables, the “`mk.pois.lk`” function returns a function of “`lambda`” that computes the likelihood function for the Poisson parameter.


$$\begin{aligned} L(\lambda) &= f_{X_1, \dots, X_n}(x_1, \dots, x_n; \lambda) \\ &= \prod_{i=1}^n f_{X_i}(x_i; \lambda) \end{aligned}$$

Using the Likelihood

Once a likelihood function is defined, it can be plotted or maximized to obtain a maximum likelihood estimate λ .

```
> PoissonData = rpois(20, 10)
> L = mk.pois.lk(PoissonData)
> x = seq(5, 15, length = 1001)
> y = sapply(x, L)
> plot(x, y, type = "l")
```

(Note that the “`sapply`” call is necessary because “`L`” is not vectorized.)



Computing a Maximum Likelihood Estimate

The likelihood function can be maximized using the “optimize” function.

```
> optimize(L, c(8, 12), maximum = TRUE)
```

```
$maximum
```

```
[1] 9.699984
```

```
$objective
```

```
[1] 7.107315e-23
```

This tells us that the maximum-likelihood estimate of λ in this case is about 9.7. Likelihood theory also provides a way of getting standard errors for the estimate (See STATS 330 for the details).

Why Are Closures Useful?

- In the last example, we could have just used a global variable “**x**” to hold the data values for the likelihood.
- However, if we need several likelihood functions, they can't all keep their data in the same “**x**.”
- Closures provide a way for functions to have their own private copies of data.
- This provides a very powerful programming tool.
- It gets used a lot in other programming languages, but not so much in R.

Alternative Way of Creating Closures

- There are a number of ways of creating closures that provide a more direct way of creating closures.
- The most direct of these are based on the “`with`” and “`local`” functions.
- They both work by directly providing a set of variables that can be closed over.
- (The use of function calls to return closures can have some rather nasty side-effects.)

Creating Closures using “with”

- The “with” function makes the elements of a named list available as variables to evaluate and expression.
- If this expression creates a function, it closes over the variables made from the list.

```
> L = with(list(x = PoissonData),  
           function(lambda)  
             prod(dpois(x, lambda)))
```

- This produces a closure that performs exactly like the earlier one created with “mk.pois.lk.”

Creating Closures using “local”

- Another direct way of creating closures is to establish local scope using the “`local`” function.
- Again, a function returned from this local scope closes over the variables defined in that scope.

```
> L = local({  
    x = PoissonData  
    function(lambda)  
      prod(dpois(x, lambda))  
  })
```

- This is the method I most commonly use for creating closures (although I do sometimes use the other methods too.)

Using Closures to Hide Helper Functions

- Often when writing software it is common to decompose a problem into smaller tasks and write functions to carry out these smaller tasks.
- Usually there is really only one function of interest, which is the one that solves the “big” problem.
- In that case there is no reason for all the other functions and associated data to be visible.
- Closures provide a perfect way of hiding things you don't want visible.

Example of Hiding Helper Functions

The following code shows how to hide two helper functions in a closure. (The closure computes binomial probabilities, directly from the formula.)

```
> binp = local({  
  fact =  
    function(n) gamma(n + 1)  
  bincoef =  
    function(n, k)  
      fact(n)/(fact(k) * fact(n - k))  
  function(k, n, p)  
    bincoef(n, k) * p^k * (1 - p)^(n - k)  
})
```

Hiding Helper Functions (Continued)

- The “`binp`” function closes over the two functions “`factor`” and “`bincoef`.”
- Those two functions are visible in the body of “`binp`,” but invisible to anything else.

```
> binp(0:5, 10, .5)
[1] 0.0009765625 0.0097656250 0.0439453125
[4] 0.1171875000 0.2050781250 0.2460937500
```

```
> objects()
[1] "binp"
```

Defining Anonymous Recursive Functions

- The factorial function can be defined recursively as follows.

```
> fact =  
  function(n)  
    if (n <= 1) 1 else n * fact(n - 1)
```

```
> fact(10)  
[1] 3628800
```

- The factorial calls itself by using its own name.
- This makes it seem like the factorial function must be assigned a fixed name (in this case “`fact`”).

An Anonymous Factorial Function

- Using a closure, it is possible to define an anonymous version of the factorial function.

```
> (local({  
    me = function(n)  
      if (n <= 1) 1 else n * me(n - 1)  
    }))(10)  
[1] 3628800
```

```
> objects()  
character(0)
```

- This factorial function has a hidden name, known only to itself, that it can use to call itself.